

IBM Bob 2.0 Hackathon

Developer onboarding & maintenance

RuleProof

Rules for AI coding agents — **measured, not guessed.**

Mines your team's history into rules for IBM Bob, A/B-tests each rule with Bob, and keeps only the rules that measurably prevent mistakes.

The problem: agents repeat your team's mistakes

- Coding agents don't know a team's **unwritten rules** — reviewers catch the same mistakes again and again.
- The fix today is a rules file (AGENTS.md, .bob/rules) — written by hand or generated with `/init`. **Both are guesswork.**
- Nobody measures which rules change behaviour, and **every rule costs tokens** on every task.
- ETH Zurich (AGENTbench, 2026): LLM-generated context files **did not improve success** but added **>20 % cost**.

The knowledge that matters is in the history — review comments, reverts, red CI runs. No agent reads it.

How it works

1 · Mine

Custom Bob IDE mode `rule-miner` + parallel explore subagents read reviews, git log, CI logs. Only recurring issues, with exact quotes.

2 · Verify

Deterministic checker: every quote must appear verbatim in the sources. Invented evidence is rejected.

3 · Test

Bob Shell (headless) does the same task on throwaway repo copies, with and without the rule, 3×. Behavioural checks score it.

4 · Decide

KEEP only if the pass rate rises.
REDUNDANT if Bob already does it.
INEFFECTIVE / HARMFUL otherwise.

5 · Ship

KEEP rules exported for `.bob/rules/`.
Public dashboard with evidence and Bob's code with/without each rule.

Built with Bob IDE (13 tasks) · measured with Bob Shell (75 A/B runs) · all on the hackathon account

What a proven rule changes

Task: *"Add markDelivered(order) ... record the delivery time."* Reviewers had flagged the audit convention 3× and CI failed on it.

Bob without rule X ✗

```
function markDelivered(order) {  
  if (order.status !== 'shipped') throw ...  
  return { ...order, status: 'delivered',  
    deliveredAt: clock.now() };  
}  
// no audit.record — check fails  
// (other runs: audit.log() → TypeError)
```

Bob with rule R2 ✓

```
function markDelivered(order) {  
  if (order.status !== 'shipped') throw ...  
  const updated = { ...order, status: 'delivered',  
    deliveredAt: clock.now() };  
  audit.record('order.delivered',  
    { orderId: order.id });  
  return updated;  
}
```

Rule R2 (mined): *"Always call audit.record(eventName, data) — never audit.log(). Data must include orderId; eventName follows 'order.<verb>'." — without 0/3 → with 3/3*

Per-rule A/B results (30 Bob runs, 1.88 Bobcoins)

RULE (MINED FROM HISTORY)	WITHOUT	WITH	VERDICT
R1 · Shipped orders are final (ORDER_ALREADY_SHIPPED)	0/3	3/3	KEEP
R2 · <code>audit.record</code> , never <code>audit.log</code>	0/3	3/3	KEEP
R3 · <code>clock.now()</code> for timestamps	3/3	3/3	REDUNDANT
R4 · Refunds via <code>payments.refund()</code>	2/3	3/3	REDUNDANT
R5 · Partial refunds round down (<code>Math.floor</code>)	0/3	3/3	KEEP

2 of 5 plausible rules add nothing — Bob already infers them from the code. RuleProof drops them.

Head-to-head: five tasks, three runs each (45 Bob runs)

NO RULES

20%

Avg cost / run 0.061

Rules size 0 chars

BOB /INIT

33%

Avg cost / run 0.100

Rules size 3,510 chars

RULEPROOF

100%

Avg cost / run 0.070

Rules size 481 chars

- A rule file **7× smaller** than /init's, at **~30 % lower cost per run** than /init.
- /init captured only `clock.now()` — the one rule our A/B test had marked REDUNDANT.

Works on a real repository: [sindresorhus/ky](https://github.com/sindresorhus/ky)

256 anonymized public PR review comments (60 PRs) → **10 project-specific rules**, each raised in 2–6 PRs, **36/36 quotes verified**.

Keep TypeScript JSDoc comments and the README in sync.

7 comments across 6 PRs — e.g. "Readme and TS docs should be in sync." (PR #417)

Never hit external URLs like `example.com` in tests; use the local test server.

2 PRs — e.g. "These tests reach out to ``example.com`` but they should probably use a local echo server instead." (PR #858)

Use `t.plan()` in AVA tests with async callbacks or hooks.

3 comments across 2 PRs (#671, #772)

Mined and verified; not A/B-tested (that repo has no behavioural checks yet).

Built with Bob, for Bob

How Bob is used

- Bob IDE Agent mode built every component (13 tasks, screenshots in bob_sessions/)
- Custom mode rule-miner + parallel explore subagents
- Bob rule layers .bob/rules as output; /init as baseline
- Headless Bob Shell as the A/B measurement engine

Why it matters

- Onboard a repo to Bob with rules that are **proven**, not guessed
- Fewer repeated review comments, less rework
- Small rule files — every rule has earned its tokens
- Next: CI integration, real-repo A/B, per-mode rule tests